Problem1

May 3, 2026

0.1 2.1 Cart-pole swing-up with limited actuation

(a)

The convex approximation (LOCP) around the current iterate $(s^{(i)}, u^{(i)})$ is

$$\begin{aligned}
\text{minimize}_{s,\,u} \quad & (s_N - s_{\text{goal}})^\top P\,(s_N - s_{\text{goal}}) \\
& + \sum_{k=0}^{N-1} \left[(s_k - s_{\text{goal}})^\top Q\,(s_k - s_{\text{goal}}) + u_k^\top R\,u_k\right] \\
\text{subject to} \quad & s_0 = \bar{s}_0 \\
& u_k \in \mathcal{U}, \quad \forall k \in \{0, \dots, N-1\} \\
& \|u_k - u_k^{(i)}\|_\infty \le \rho, \quad \forall k \in \{0, \dots, N-1\} \\
& \|s_k - s_k^{(i)}\|_\infty \le \rho, \quad \forall k \in \{0, \dots, N\} \\
& s_{k+1} = A_k s_k + B_k u_k + c_k, \quad \forall k \in \{0, \dots, N-1\}
\end{aligned}$$

where the terminal constraint $s_N = s_{\text{goal}}$ is replaced by the terminal cost $(s_N - s_{\text{goal}})^\top P\,(s_N - s_{\text{goal}})$, and A_k, B_k, c_k come from the linearization as follow:

$$s_{k+1} = f_{\mathbf{d}}\big(s_k^{(i)}, u_k^{(i)}\big) + \frac{\partial f_{\mathbf{d}}}{\partial s}\big(s_k^{(i)}, u_k^{(i)}\big)\big(s_k - s_k^{(i)}\big) + \frac{\partial f_{\mathbf{d}}}{\partial u}\big(s_k^{(i)}, u_k^{(i)}\big)\big(u_k - u_k^{(i)}\big) = A_k s_k + B_k u_k + c_k$$

with

$$A_k = \frac{\partial f_{\mathbf{d}}}{\partial s}\big(s_k^{(i)}, u_k^{(i)}\big), \qquad B_k = \frac{\partial f_{\mathbf{d}}}{\partial u}\big(s_k^{(i)}, u_k^{(i)}\big), \qquad c_k = f_{\mathbf{d}}\big(s_k^{(i)}, u_k^{(i)}\big) - A_k s_k^{(i)} - B_k u_k^{(i)}$$

(b) Please see `cartpole_swingup_constrained.py`.

(c)

Please see `cartpole_swingup_constrained.py`.

(d)

[1]:
```
%run cartpole_swingup_constrained.py
```

```
 88%|      | 88/100 [01:12<00:09,  1.22it/s, objective change=0.32265]
```

```
SCP converged after 88 iterations.
```

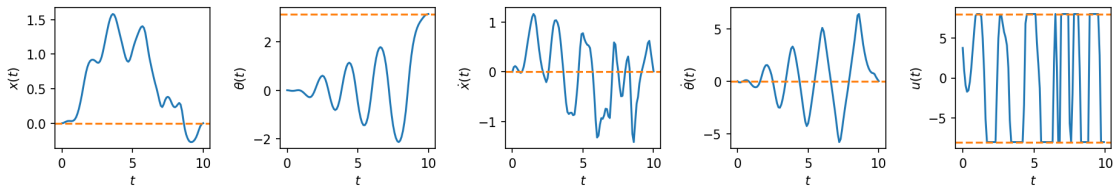

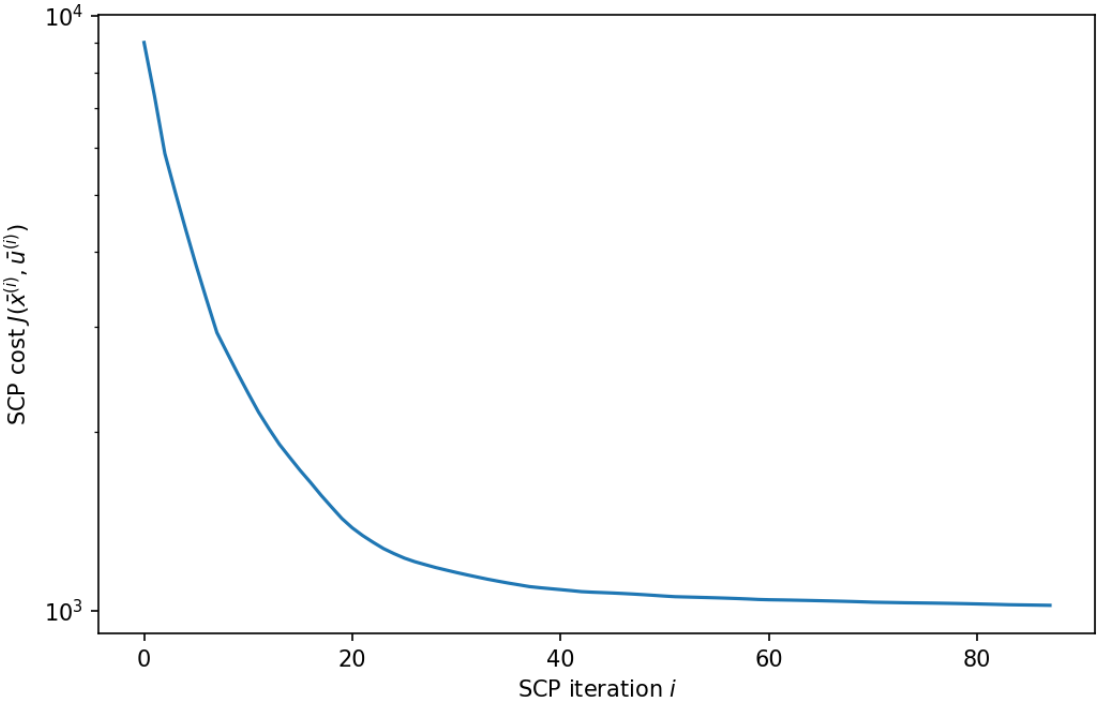

The pendulum just reaches s_{goal} at the end of the horizon but does not stay there. A natural scheme is to use SCP to generate the swing-up trajectory, then switch to a closed-loop LQR controller linearized around the upright equilibrium to keep the pendulum there indefinitely.

“““ Starter code for the problem “Cart-pole swing-up with limited actuation”.

Autonomous Systems Lab (ASL), Stanford University “““

from functools import partial

from animations import animate_cartpole

import cvxpy as cvx

import jax import jax.numpy as jnp

import matplotlib.pyplot as plt

import numpy as np

from tqdm import tqdm

@partial(jax.jit, static_argnums=(0,)) @partial(jax.vmap, in_axes=(None, 0, 0)) def affinize(f, s, u): “““Affinize the function `f(s, u)` around `(s, u)`.

```
Arguments
---------
f : callable
    A nonlinear function with call signature `f(s, u)`.
s : numpy.ndarray
    The state (1-D).
u : numpy.ndarray
    The control input (1-D).

Returns
-------
A : jax.numpy.ndarray
    The Jacobian of `f` at `(s, u)`, with respect to `s`.
B : jax.numpy.ndarray
    The Jacobian of `f` at `(s, u)`, with respect to `u`.
c : jax.numpy.ndarray
    The offset term in the first-order Taylor expansion of `f` at `(s, u)`
    that sums all vector terms strictly dependent on the nominal point
    `(s, u)` alone.
"""
# PART (b) #############################################################
# INSTRUCTIONS: Use JAX to affinize `f` around `(s, u)` in two lines.
A,B=jax.jacobian(f, argnums=(0,1))(s,u)
c=f(s,u)-A@s-B@u
# END PART (b) #############################################################
return A, B, c
```

def solve_swingup_scp(f, s0, s_goal, N, P, Q, R, u_max, , eps, max_iters): “““Solve the cart-pole swing-up problem via SCP.

```
Arguments
---------
f : callable
    A function describing the discrete-time dynamics, such that
    `s[k+1] = f(s[k], u[k])`.
s0 : numpy.ndarray
    The initial state (1-D).
s_goal : numpy.ndarray
    The goal state (1-D).
```

```
N : int
    The time horizon of the LQR cost function.
P : numpy.ndarray
    The terminal state cost matrix (2-D).
Q : numpy.ndarray
    The state stage cost matrix (2-D).
R : numpy.ndarray
    The control stage cost matrix (2-D).
u_max : float
    The bound defining the control set `[-u_max, u_max]`.
ρ : float
    Trust region radius.
eps : float
    Termination threshold for SCP.
max_iters : int
    Maximum number of SCP iterations.

Returns
-------
s : numpy.ndarray
    A 2-D array where `s[k]` is the open-loop state at time step `k`,
    for `k = 0, 1, ..., N-1`
u : numpy.ndarray
    A 2-D array where `u[k]` is the open-loop state at time step `k`,
    for `k = 0, 1, ..., N-1`
J : numpy.ndarray
    A 1-D array where `J[i]` is the SCP sub-problem cost after the i-th
    iteration, for `i = 0, 1, ..., (iteration when convergence occured)`
"""
n = Q.shape[0]  # state dimension
m = R.shape[0]  # control dimension

# Initialize dynamically feasible nominal trajectories
u = np.zeros((N, m))
s = np.zeros((N + 1, n))
s[0] = s0
for k in range(N):
    s[k + 1] = fd(s[k], u[k])

# Do SCP until convergence or maximum number of iterations is reached
converged = False
J = np.zeros(max_iters + 1)
J[0] = np.inf
for i in (prog_bar := tqdm(range(max_iters))):
    s, u, J[i + 1] = scp_iteration(f, s0, s_goal, s, u, N, P, Q, R, u_max, ρ)
    dJ = np.abs(J[i + 1] - J[i])
    prog_bar.set_postfix({"objective change": "{:.5f}".format(dJ)})
    if dJ < eps:
        converged = True
        print("SCP converged after {} iterations.".format(i))
        break
if not converged:
    raise RuntimeError("SCP did not converge!")
J = J[1 : i + 1]
```

```
return s, u, J
```

def scp_iteration(f, s0, s_goal, s_prev, u_prev, N, P, Q, R, u_max,): """Solve a single SCP sub-problem for the cart-pole swing-up problem.

```
Arguments
---------
f : callable
    A function describing the discrete-time dynamics, such that
    `s[k+1] = f(s[k], u[k])`.
s0 : numpy.ndarray
    The initial state (1-D).
s_goal : numpy.ndarray
    The goal state (1-D).
s_prev : numpy.ndarray
    The state trajectory around which the problem is convexified (2-D).
u_prev : numpy.ndarray
    The control trajectory around which the problem is convexified (2-D).
N : int
    The time horizon of the LQR cost function.
P : numpy.ndarray
    The terminal state cost matrix (2-D).
Q : numpy.ndarray
    The state stage cost matrix (2-D).
R : numpy.ndarray
    The control stage cost matrix (2-D).
u_max : float
    The bound defining the control set `[-u_max, u_max]`.
ρ : float
    Trust region radius.

Returns
-------
s : numpy.ndarray
    A 2-D array where `s[k]` is the open-loop state at time step `k`,
    for `k = 0, 1, ..., N-1`
u : numpy.ndarray
    A 2-D array where `u[k]` is the open-loop state at time step `k`,
    for `k = 0, 1, ..., N-1`
J : float
    The SCP sub-problem cost.
"""
A, B, c = affinize(f, s_prev[:-1], u_prev)
A, B, c = np.array(A), np.array(B), np.array(c)
n = Q.shape[0]
m = R.shape[0]
s_cvx = cvx.Variable((N + 1, n))
u_cvx = cvx.Variable((N, m))

# PART (c) #############################################################
# INSTRUCTIONS: Construct the convex SCP sub-problem.
objective = 0.0
constraints = []
constraints += [s_cvx[0] == s0]
constraints+=[cvx.norm(s_cvx[k]-s_prev[k],'inf')<=ρ for k in range(N+1)]
```

```
constraints+=[cvx.norm(u_cvx[k]-u_prev[k],'inf')<=ρ for k in range(N)]
constraints+=[cvx.norm(u_cvx[k], 'inf') <= u_max for k in range(N)]
constraints += [s_cvx[k + 1] == A[k] @ s_cvx[k] + B[k] @ u_cvx[k] + c[k] for k in range(N)]
objective += cvx.quad_form(s_cvx[N] - s_goal, P)
objective += cvx.sum([cvx.quad_form(s_cvx[k] - s_goal, Q) + cvx.quad_form(u_cvx[k], R) for k in range(N)])
# END PART (c) ############################################################

prob = cvx.Problem(cvx.Minimize(objective), constraints)
prob.solve()
if prob.status != "optimal":
    raise RuntimeError("SCP solve failed. Problem status: " + prob.status)
s = s_cvx.value
u = u_cvx.value
J = prob.objective.value
return s, u, J
```

def cartpole(s, u): “” “Compute the cart-pole state derivative.” “” mp = 1.0 # pendulum mass mc = 4.0 # cart mass L = 1.0 # pendulum length g = 9.81 # gravitational acceleration

```
x, θ, dx, dθ = s
sinθ, cosθ = jnp.sin(θ), jnp.cos(θ)
h = mc + mp * (sinθ**2)
ds = jnp.array(
    [
        dx,
        dθ,
        (mp * sinθ * (L * (dθ**2) + g * cosθ) + u[0]) / h,
        -((mc + mp) * g * sinθ + mp * L * (dθ**2) * sinθ * cosθ + u[0] * cosθ)
        / (h * L),
    ]
)
return ds
```

def discretize(f, dt): “” “Discretize continuous-time dynamics `f` via Runge-Kutta integration.” “”

```
def integrator(s, u, dt=dt):
    k1 = dt * f(s, u)
    k2 = dt * f(s + k1 / 2, u)
    k3 = dt * f(s + k2 / 2, u)
    k4 = dt * f(s + k3, u)
    return s + (k1 + 2 * k2 + 2 * k3 + k4) / 6

return integrator
```

Define constants

n = 4 # state dimension m = 1 # control dimension s_goal = np.array([0, np.pi, 0, 0]) # desired upright pendulum state s0 = np.array([0, 0, 0, 0]) # initial downright pendulum state dt = 0.1 # discrete time resolution T = 10.0 # total simulation time P = 1e3 * np.eye(n) # terminal state cost matrix Q = np.diag([1e-2, 1.0, 1e-3, 1e-3]) # state cost matrix R = 1e-3 * np.eye(m) # control cost matrix = 1.0 # trust region parameter u_max = 8.0 # control effort bound eps = 5e-1 # convergence tolerance max_iters = 100 # maximum number of SCP iterations animate = False # flag for animation

Initialize the discrete-time dynamics

fd = jax.jit(discretize(cartpole, dt))

Solve the swing-up problem with SCP

t = np.arange(0.0, T + dt, dt) N = t.size - 1 s, u, J = solve_swingup_scp(fd, s0, s_goal, N, P, Q, R, u_max, , eps, max_iters)

Simulate open-loop control

for k in range(N): s[k + 1] = fd(s[k], u[k])

Plot state and control trajectories

fig, ax = plt.subplots(1, n + m, dpi=150, figsize=(15, 2)) plt.subplots_adjust(wspace=0.45) labels_s = (r"$x(t)$“, r"$\theta(t)$“, r"$\dot{x}(t)$“, r"$\dot{\theta}(t)$“) labels_u = (r"$u(t)$“,) for i in range(n): ax[i].plot(t, s[:, i]) ax[i].axhline(s_goal[i], linestyle="–“, color="tab:orange") ax[i].set_xlabel(r"t“) ax[i].set_ylabel(labels_s[i]) for i in range(m): ax[n + i].plot(t[:-1], u[:, i]) ax[n + i].axhline(u_max, linestyle="–“, color="tab:orange") ax[n + i].axhline(-u_max, linestyle="–“, color="tab:orange") ax[n + i].set_xlabel(r"t“) ax[n + i].set_ylabel(labels_u[i]) plt.savefig("cartpole_swingup_constrained.png", bbox_inches="tight")

Plot cost history over SCP iterations

fig, ax = plt.subplots(1, 1, dpi=150, figsize=(8, 5)) ax.semilogy(J) ax.set_xlabel(r"SCP iteration i“) ax.set_ylabel(r"SCP cost $J(\bar{x}^{(i)}, \bar{u}^{(i)})$“) plt.savefig("cartpole_swingup_constrained_cost.png", bbox_inches="tight") plt.show()

Animate the solution

if animate: fig, ani = animate_cartpole(t, s[:, 0], s[:, 1]) ani.save(“cartpole_swingup_constrained.mp4”, writer=“ffmpeg”) plt.show()

Problem2

May 3, 2026

0.1 2.2 Shortest path through a grid

(a)

Doing DP backward (B to A) we compute at each node the optimal cost-to-go (J^* in the figure) and store the optimal edge to follow, represented as a blue arrow. Finally, the optimal path from A to B is computed by always choosing the optimal edge and is highlighted in red. It corresponds to: **up, up, down, up, down, down**.

[1]:
```
from IPython.display import Image, display
display(Image(filename='/Users/erwinpoussi/Downloads/NBMetadataCache 2.png',␣
↪width=700))
```

ight and t tr
ecision to go node. Given
ow-dimensional state space, we will use dynamic prog
each node in the grid.

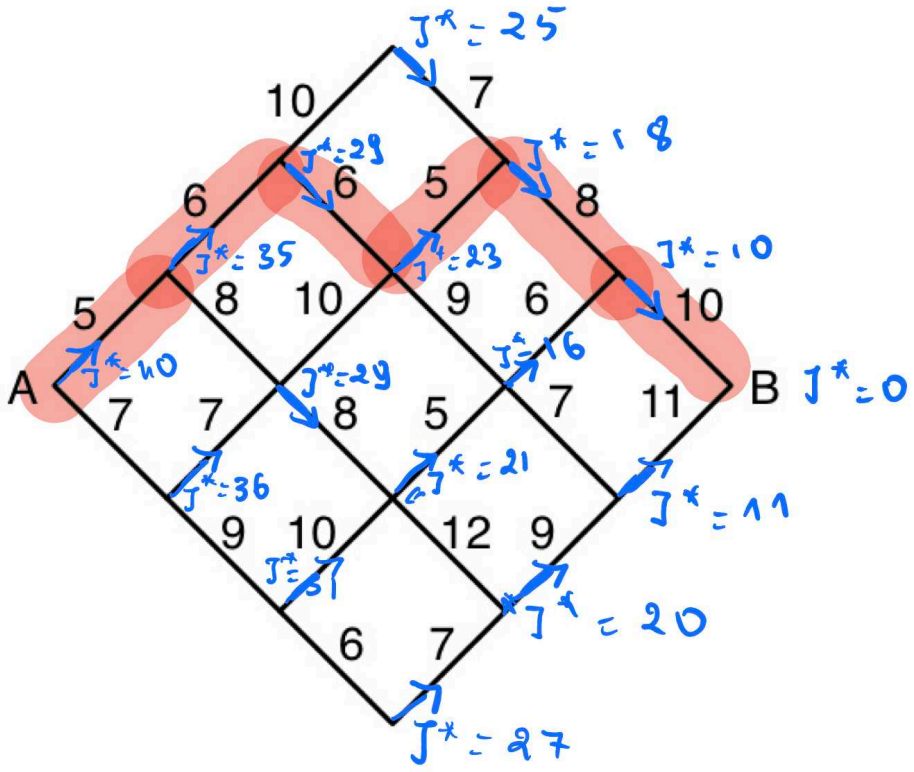

Figure 1: Shortest path problem on a grid.

amming (DP) to find the shortest path from A to B
lized version of the shortest path problem in Figure
h side. Find the number of computations required b
number of routes that such an algorithm would nee

(b)

Number of nodes. For a grid with n segments per side, the diamond has columns $0, 1, \ldots, 2n$,

with the number of nodes per column growing from 1 to $n+1$ then shrinking back to 1:

$$n_{\text{nodes}} = \sum_{k=0}^{n}(k+1) + \sum_{k=1}^{n} k = \frac{(n+1)(n+2)}{2} + \frac{n(n+1)}{2} = (n+1)^2.$$

For $n = 3$: $n_{\text{nodes}} = 16$.

Exhaustive search. Every path from A to B consists of $2n$ steps, each either up or down, and must end at B (same row as A), which requires exactly n ups and n downs. The number of such paths is therefore

$$\boxed{N_{\text{exhaustive}} = \binom{2n}{n} = \frac{(2n)!}{(n!)^2}}.$$

For $n = 3$: $\binom{6}{3} = 20$ routes.

Dynamic programming. The DP backward sweep evaluates the cost-to-go J^* once at every node except the terminal B. At each node the computation is a single min over at most 2 successors, so the total number of DP evaluations equals the number of non-terminal nodes:

$$\boxed{N_{\text{DP}} = (n+1)^2 - 1 = n^2 + 2n}.$$

For $n = 3$: $N_{\text{DP}} = 9 + 6 = 15$ evaluations.

Comparison. $\binom{2n}{n}$ grows exponentially ($\sim 4^n/\sqrt{\pi n}$) while $n^2 + 2n$ grows only polynomially, illustrating the considerable computational advantage of DP over exhaustive search.

problem3

May 3, 2026

0.1 2.3 Cart-pole balance

(a)

Starting from the continuous-time dynamics $\dot{s} = f(s, u)$, we apply **Euler integration** to get the discrete-time approximation:

$$s_{k+1} \approx s_k + \Delta t\, f(s_k, u_k).$$

We then **linearize** f around the equilibrium $(\bar{s}, \bar{u})$ via a first-order Taylor expansion:

$$f(s_k, u_k) \approx f(\bar{s}, \bar{u}) + \left.\frac{\partial f}{\partial s}\right|_{(\bar{s},\bar{u})} (s_k - \bar{s}) + \left.\frac{\partial f}{\partial u}\right|_{(\bar{s},\bar{u})} (u_k - \bar{u}).$$

Since $(\bar{s}, \bar{u})$ is an equilibrium, $f(\bar{s}, \bar{u}) = 0$, and $\bar{u} = 0$. Defining $\tilde{s}_k = s_k - \bar{s}$ and subtracting $\bar{s}$ from both sides:

$$\tilde{s}_{k+1} \approx \tilde{s}_k + \Delta t \left(\left.\frac{\partial f}{\partial s}\right|_{(\bar{s},\bar{u})} \tilde{s}_k + \left.\frac{\partial f}{\partial u}\right|_{(\bar{s},\bar{u})} u_k \right) = \underbrace{\left(I + \Delta t \left.\frac{\partial f}{\partial s}\right|_{(\bar{s},\bar{u})} \right)}_{A} \tilde{s}_k + \underbrace{\Delta t \left.\frac{\partial f}{\partial u}\right|_{(\bar{s},\bar{u})}}_{B} u_k.$$

Substituting the given Jacobians evaluated at $\bar{s} = (0, \pi, 0, 0)$, $\bar{u} = 0$:

$$A = I_4 + \Delta t \begin{bmatrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & \frac{m_p g}{m_c} & 0 & 0 \\ 0 & \frac{(m_c+m_p)g}{m_c \ell} & 0 & 0 \end{bmatrix}, \qquad B = \Delta t \begin{bmatrix} 0 \\ 0 \\ \frac{1}{m_c} \\ \frac{1}{m_c \ell} \end{bmatrix}.$$

(b)

[4]:
```
import numpy as np

mp=2
mc=10
l=1
g=9.81
dt=0.1
```

```
Q=np.eye(4)
R=np.eye(1)
A=np.eye(4)+dt*np.array([[0,0,1,0],[0,0,0,1],[0,mp*g/mc,0,0],[0,(mc+mp)*g/
  ↪(mc*l),0,0]])
B=np.array([[0],[0],[dt/mc],[dt/(mc*l)]])

P_inf=np.zeros((4,4))
K_inf=-np.linalg.inv((R+B.T@P_inf@B))@B.T@P_inf@A
Pnew=Q+A.T@P_inf@(A+B@K_inf)
while np.max(np.abs(Pnew-P_inf))>1e-4:
    P_inf=Pnew
    K_inf=-np.linalg.inv(R+B.T@P_inf@B)@B.T@P_inf@A
    Pnew=Q+A.T@P_inf@(A+B@K_inf)
print("P_inf:\n",np.round(P_inf,2))
print("K_inf:\n",np.round(K_inf,2))
```

```
P_inf:
 [[ 5.787000e+01 -9.360000e+02  1.595600e+02 -2.967100e+02]
 [-9.360000e+02  1.259655e+05 -5.026500e+03  3.750146e+04]
 [ 1.595600e+02 -5.026500e+03  8.120600e+02 -1.592050e+03]
 [-2.967100e+02  3.750146e+04 -1.592050e+03  1.118159e+04]]
K_inf:
 [[   0.73 -231.85    4.22  -68.25]]
```

(c)

[11]:
```
from scipy.integrate import odeint
from cartpole_balance import cartpole
import matplotlib.pyplot as plt

def cartpole_odeint(s, t, u):
    return cartpole(s, u)

s_bar = np.array([0, np.pi, 0, 0])
s0 = np.array([0, 3*np.pi/4, 0, 0])
T = 30
N = int(T/dt)
t = np.arange(0, T+dt, dt)

s = np.zeros((N+1, 4))
u_hist = np.zeros(N)
s[0] = s0

for k in range(N):
    u_hist[k] = (K_inf @ (s[k] - s_bar))[0]
    s[k+1] = odeint(cartpole_odeint, s[k], t[k:k+2], (np.
  ↪array([u_hist[k]]),))[1]
```

```
fig, axes = plt.subplots(5, 1, figsize=(8, 10), sharex=True)
labels = ['x', ' ', 'ẋ', '̇', 'u']
refs = [s_bar[0], s_bar[1], s_bar[2], s_bar[3], 0]
for i, ax in enumerate(axes):
    ax.plot(t if i < 4 else t[:-1], s[:, i] if i < 4 else u_hist, label='state'␣
 ↪if i < 4 else 'u')
    ax.axhline(refs[i], linestyle='--', color='orange', label='reference')
    ax.set_ylabel(labels[i])
    ax.legend(fontsize=7)
axes[-1].set_xlabel('t (s)')
plt.tight_layout()
plt.show()
```

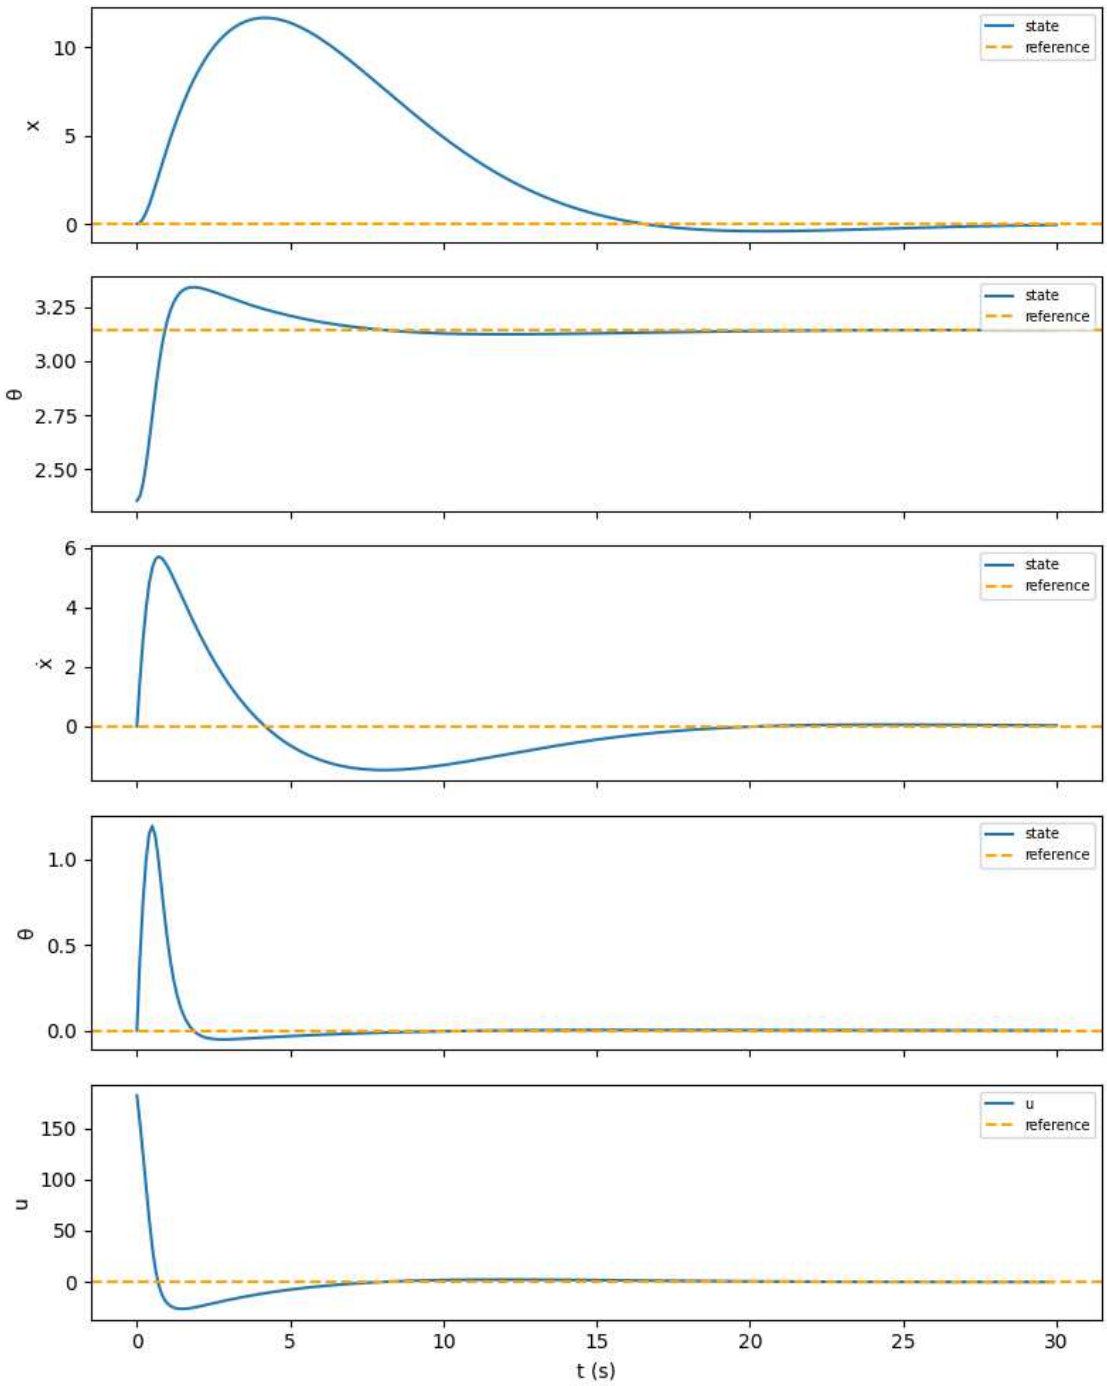

[ ]:
```
# Animation
from animations import animate_cartpole
from IPython.display import HTML

fig, ani = animate_cartpole(t, s[:,0], s[:,1])
```

```
plt.close(fig)
HTML(ani.to_jshtml())
```

[ ]: <IPython.core.display.HTML object>

(d)

(i) The Jacobians $\frac{\partial f}{\partial s}$ and $\frac{\partial f}{\partial u}$ do not depend on the cart position x — only on θ, $\dot{\theta}$, and u. Since the reference trajectory always maintains $\bar{\theta}(t) = \pi$, $\dot{\bar{\theta}}(t) = 0$, and $\bar{u} = 0$, the linearization point is identical at every time step regardless of the time-varying $\bar{x}(t)$. Therefore A and B are constant and do not need to be recomputed.

(iii) The desired trajectory $\bar{x}(t) = a\sin(2\pi t/T)$ requires the cart to accelerate ($\ddot{\bar{x}} \neq 0$), which physically means a nonzero reference force $\bar{u}(t) \neq 0$ is needed to drive it. The LQR was designed assuming $\bar{u} = 0$, so there is no feedforward term to generate this force. As a result the controller cannot simultaneously track $\bar{x}(t)$ and keep $\theta = \pi$: the cart acceleration perturbs the pendulum, causing θ and $\dot{\theta}$ to oscillate. Increasing Q tightens regulation around the reference but does not fix the missing feedforward, so θ and $\dot{\theta}$ keep oscillating.

[8]:
```
from scipy.integrate import odeint
from cartpole_balance import cartpole

def cartpole_odeint(s, t, u):
    return cartpole(s, u)

T = 30
N = int(T/dt)
t = np.arange(0, T+dt, dt)

a, T_ref = 10, 10

def s_bar_tv(t_val):
    return np.array([a*np.sin(2*np.pi*t_val/T_ref), np.pi, a*(2*np.pi/T_ref)*np.
 ↪cos(2*np.pi*t_val/T_ref), 0])

s0_d = np.array([0, np.pi, 0, 0])
s_d = np.zeros((N+1, 4))
u_d = np.zeros(N)
s_d[0] = s0_d

for k in range(N):
    u_d[k] = (K_inf @ (s_d[k] - s_bar_tv(t[k])))[0]
    s_d[k+1] = odeint(cartpole_odeint, s_d[k], t[k:k+2], (np.
 ↪array([u_d[k]]),))[1]

s_ref_hist = np.array([s_bar_tv(tk) for tk in t])

fig, axes = plt.subplots(5, 1, figsize=(8, 10), sharex=True)
```

```
labels = ['x', ' ', 'ẋ', '¨', 'u']
for i, ax in enumerate(axes):
    if i < 4:
        ax.plot(t, s_d[:, i], label='state')
        ax.plot(t, s_ref_hist[:, i], '--', label='reference')
        ax.legend(fontsize=7)
    else:
        ax.plot(t[:-1], u_d)
    ax.set_ylabel(labels[i])
axes[-1].set_xlabel('t (s)')
plt.tight_layout()
plt.show()
```

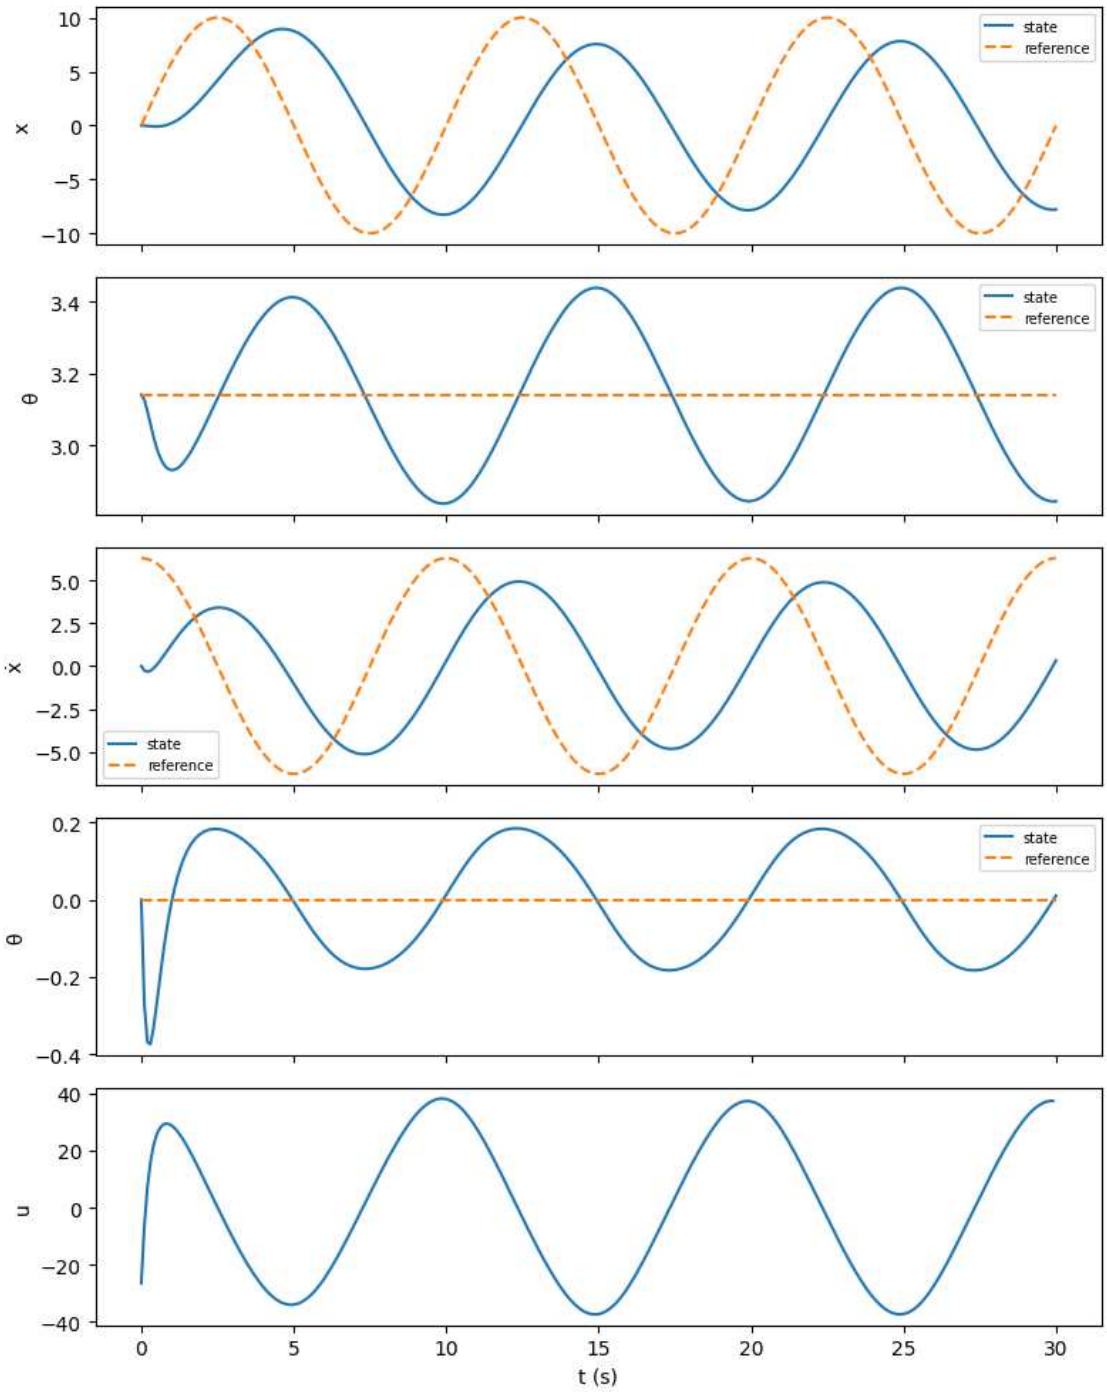

[9]:
```
# Animation
from animations import animate_cartpole
from IPython.display import HTML

fig, ani = animate_cartpole(t, s_d[:,0], s_d[:,1])
```

```
plt.close(fig)
HTML(ani.to_jshtml())
```

[9]: <IPython.core.display.HTML object>

[ ]:

problem4

May 3, 2026

0.1 2.4 Cart-pole swing-up

(a) Please see `cartpole_swingup.py`.

(b)

Using a second order Taylor expansion of the stage cost $c(\tilde{s}_k, \tilde{u}_k)$ around $(\bar{s}_k, \bar{u}_k)$:

$$c(\tilde{s}_k, \tilde{u}_k) = c(\bar{s}_k, \bar{u}_k) + \underbrace{\left(\frac{\partial c}{\partial s}\bigg|_{(\bar{s}_k, \bar{u}_k)} \right)^\top}_{=q_k} \tilde{s}_k + \underbrace{\left(\frac{\partial c}{\partial u}\bigg|_{(\bar{s}_k, \bar{u}_k)} \right)^\top}_{=r_k} \tilde{u}_k + \frac{1}{2}\tilde{s}_k^\top \underbrace{\frac{\partial^2 c}{\partial s^2}\bigg|_{(\bar{s}_k, \bar{u}_k)}}_{=Q} \tilde{s}_k + \frac{1}{2}\tilde{u}_k^\top \underbrace{\frac{\partial^2 c}{\partial u^2}\bigg|_{(\bar{s}_k, \bar{u}_k)}}_{=R} \tilde{u}_k$$

hence $\boxed{q_k = Q(\bar{s}_k - s_{\text{goal}})}$ and $\boxed{r_k = R\bar{u}_k}$.

Same for the terminal cost $c_N(\tilde{s}_N)$:

$$c_N(\tilde{s}_N) = c_N(\bar{s}_N) + \underbrace{\left(\frac{\partial c_N}{\partial s}\bigg|_{\bar{s}_N} \right)^\top}_{=q_N} \tilde{s}_N + \frac{1}{2}\tilde{s}_N^\top \underbrace{\frac{\partial^2 c_N}{\partial s^2}\bigg|_{\bar{s}_N}}_{=Q_N} \tilde{s}_N$$

hence $\boxed{q_N = Q_N(\bar{s}_N - s_{\text{goal}})}$.

(c) Please see `cartpole_swingup.py`.

[1]:
```
%run cartpole_swingup.py
```

```
---------------------------------------------------------------------------
ModuleNotFoundError                       Traceback (most recent call last)
File ~/Documents/school projects/Optimal-learned-based-control/Hw2/notebooks/
 ↪cartpole_swingup.py:11
      7 import time
      9 from animations import animate_cartpole
---> 11 import jax
     12 import jax.numpy as jnp
     14 import matplotlib.pyplot as plt
```

```
ModuleNotFoundError: No module named 'jax'
```

(d) Please see `cartpole_swingup.py`.

[2]:
```
from IPython.display import Image, display
print("Open-loop:")
display(Image(filename='cartpole_swingup_ol.png'))
```

```
Open-loop:
```

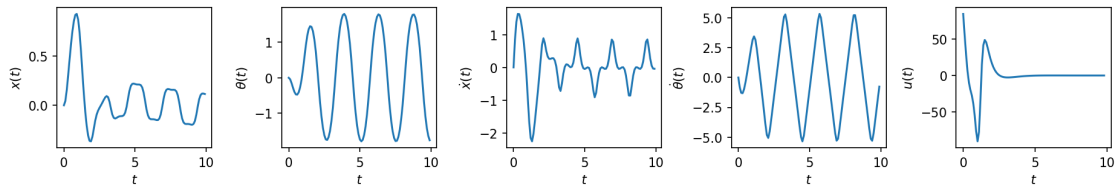

[3]:
```
print("Closed-loop:")
display(Image(filename='cartpole_swingup_cl.png'))
```

```
Closed-loop:
```

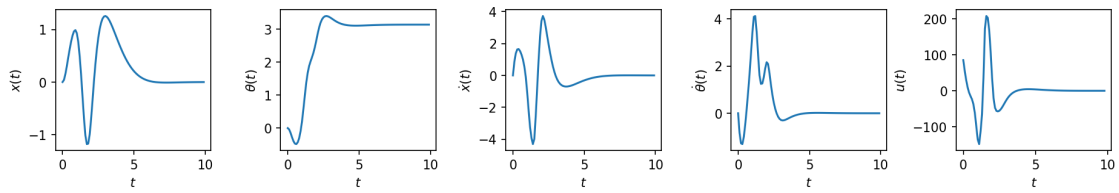

While the open-loop is consistently oscillating, the closed-loop is progressively damped towards convergence as expected.

“““ Starter code for the problem “Cart-pole swing-up”.

Autonomous Systems Lab (ASL), Stanford University “““

import time

from animations import animate_cartpole

import jax import jax.numpy as jnp

import matplotlib.pyplot as plt

import numpy as np

from scipy.integrate import odeint

def linearize(f, s, u): “““Linearize the function `f(s, u)` around `(s, u)`.

```
Arguments
---------
f : callable
    A nonlinear function with call signature `f(s, u)`.
s : numpy.ndarray
    The state (1-D).
u : numpy.ndarray
    The control input (1-D).

Returns
-------
A : numpy.ndarray
    The Jacobian of `f` at `(s, u)`, with respect to `s`.
B : numpy.ndarray
    The Jacobian of `f` at `(s, u)`, with respect to `u`.
"""
# WRITE YOUR CODE BELOW ###################################################
# INSTRUCTIONS: Use JAX to compute `A` and `B` in one line.
A, B = jax.jacobian(f, argnums=(0, 1))(s, u)
###########################################################################
return A, B
```

def ilqr(f, s0, s_goal, N, Q, R, QN, eps=1e-3, max_iters=1000): “““Compute the iLQR set-point tracking solution.

```
Arguments
---------
f : callable
    A function describing the discrete-time dynamics, such that
    `s[k+1] = f(s[k], u[k])`.
s0 : numpy.ndarray
    The initial state (1-D).
s_goal : numpy.ndarray
    The goal state (1-D).
N : int
    The time horizon of the LQR cost function.
Q : numpy.ndarray
    The state cost matrix (2-D).
R : numpy.ndarray
    The control cost matrix (2-D).
QN : numpy.ndarray
```

```
    The terminal state cost matrix (2-D).
eps : float, optional
    Termination threshold for iLQR.
max_iters : int, optional
    Maximum number of iLQR iterations.

Returns
-------
s_bar : numpy.ndarray
    A 2-D array where `s_bar[k]` is the nominal state at time step `k`,
    for `k = 0, 1, ..., N-1`
u_bar : numpy.ndarray
    A 2-D array where `u_bar[k]` is the nominal control at time step `k`,
    for `k = 0, 1, ..., N-1`
Y : numpy.ndarray
    A 3-D array where `Y[k]` is the matrix gain term of the iLQR control
    law at time step `k`, for `k = 0, 1, ..., N-1`
y : numpy.ndarray
    A 2-D array where `y[k]` is the offset term of the iLQR control law
    at time step `k`, for `k = 0, 1, ..., N-1`
"""
if max_iters <= 1:
    raise ValueError("Argument `max_iters` must be at least 1.")
n = Q.shape[0]  # state dimension
m = R.shape[0]  # control dimension

# Initialize gains `Y` and offsets `y` for the policy
Y = np.zeros((N, m, n))
y = np.zeros((N, m))

# Initialize the nominal trajectory `(s_bar, u_bar`), and the
# deviations `(ds, du)`
u_bar = np.zeros((N, m))
s_bar = np.zeros((N + 1, n))
s_bar[0] = s0
for k in range(N):
    s_bar[k + 1] = f(s_bar[k], u_bar[k])
ds = np.zeros((N + 1, n))
du = np.zeros((N, m))

# iLQR loop
converged = False
for _ in range(max_iters):
    # Linearize the dynamics at each step `k` of `(s_bar, u_bar)`
    A, B = jax.vmap(linearize, in_axes=(None, 0, 0))(f, s_bar[:-1], u_bar)
    A, B = np.array(A), np.array(B)

    # PART (c) ###############################################################
    # INSTRUCTIONS: Update `Y`, `y`, `ds`, `du`, `s_bar`, and `u_bar`.

    # Backward pass: Riccati recursion for P (value Hessian), p (value gradient),
    # gains Y[k] and offsets y[k]
    P = QN
    p = QN @ (s_bar[-1] - s_goal)
```

```
    for k in range(N - 1, -1, -1):
        inv = np.linalg.inv(R + B[k].T @ P @ B[k])
        Y[k] = -inv @ B[k].T @ P @ A[k]
        y[k] = -inv @ (R @ u_bar[k] + B[k].T @ p)
        p = Q @ (s_bar[k] - s_goal) + A[k].T @ (p + P @ B[k] @ y[k])
        P = Q + A[k].T @ P @ (A[k] + B[k] @ Y[k])

    # Forward pass: apply policy on real dynamics
    s_new = np.zeros_like(s_bar)
    u_new = np.zeros_like(u_bar)
    s_new[0] = s0
    for k in range(N):
        du[k] = Y[k] @ (s_new[k] - s_bar[k]) + y[k]
        u_new[k] = u_bar[k] + du[k]
        s_new[k + 1] = f(s_new[k], u_new[k])
    s_bar = s_new
    u_bar = u_new

    ############################################################################

    if np.max(np.abs(du)) < eps:
        converged = True
        break
if not converged:
    raise RuntimeError("iLQR did not converge!")
return s_bar, u_bar, Y, y
```

def cartpole(s, u): “‘ “‘Compute the cart-pole state derivative.”‘ ”’ mp = 2.0 # pendulum mass mc = 10.0 # cart mass L = 1.0 # pendulum length g = 9.81 # gravitational acceleration

```
x, θ, dx, dθ = s
sinθ, cosθ = jnp.sin(θ), jnp.cos(θ)
h = mc + mp * (sinθ**2)
ds = jnp.array(
    [
        dx,
        dθ,
        (mp * sinθ * (L * (dθ**2) + g * cosθ) + u[0]) / h,
        -((mc + mp) * g * sinθ + mp * L * (dθ**2) * sinθ * cosθ + u[0] * cosθ)
        / (h * L),
    ]
)
return ds
```

Define constants

n = 4 # state dimension m = 1 # control dimension Q = np.diag(np.array([10.0, 10.0, 2.0, 2.0])) # state cost matrix R = 1e-2 * np.eye(m) # control cost matrix QN = 1e2 * np.eye(n) # terminal state cost matrix s0 = np.array([0.0, 0.0, 0.0, 0.0]) # initial state s_goal = np.array([0.0, np.pi, 0.0, 0.0]) # goal state T = 10.0 # simulation time dt = 0.1 # sampling time animate = False # flag for animation closed_loop = True # flag for closed-loop control

Initialize continuous-time and discretized dynamics

f = jax.jit(cartpole) fd = jax.jit(lambda s, u, dt=dt: s + dt * f(s, u))

Compute the iLQR solution with the discretized dynamics

print(“Computing iLQR solution …”, end=“ “, flush=True) start = time.time() t = np.arange(0.0, T, dt) N = t.size - 1 s_bar, u_bar, Y, y = ilqr(fd, s0, s_goal, N, Q, R, QN) print(”done! ({:.2f} s)“.format(time.time() - start), flush=True)

Simulate on the true continuous-time system

print(“Simulating …”, end=“ “, flush=True) start = time.time() s = np.zeros((N + 1, n)) u = np.zeros((N, m)) s[0] = s0 for k in range(N): # PART (d) ## # INSTRUCTIONS: Compute either the closed-loop or open-loop value of # `u[k]`, depending on the Boolean flag `closed_loop`. if closed_loop: u[k] = u_bar[k] + Y[k] @ (s[k] - s_bar[k]) + y[k] else: # do open-loop control u[k] = u_bar[k] ## s[k + 1] = odeint(lambda s, t: f(s, u[k]), s[k], t[k : k + 2])[1] print(”done! ({:.2f} s)“.format(time.time() - start), flush=True)

Plot

fig, axes = plt.subplots(1, n + m, dpi=150, figsize=(15, 2)) plt.subplots_adjust(wspace=0.45) labels_s = (r”$x(t)$“, r”$\theta(t)$“, r”$\dot{x}(t)$“, r”$\dot{\theta}(t)$“) labels_u = (r”$u(t)$“,) for i in range(n): axes[i].plot(t, s[:, i]) axes[i].set_xlabel(r”t“) axes[i].set_ylabel(labels_s[i]) for i in range(m): axes[n + i].plot(t[:-1], u[:, i]) axes[n + i].set_xlabel(r”t“) axes[n + i].set_ylabel(labels_u[i]) if closed_loop: plt.savefig(”cartpole_swingup_cl.png”, bbox_inches=“tight”) else: plt.savefig(“cartpole_swingup_ol.png”, bbox_inches=“tight”) plt.show()

if animate: fig, ani = animate_cartpole(t, s[:, 0], s[:, 1]) ani.save(“cartpole_swingup.mp4”, writer=“ffmpeg”) plt.show()

Problem5

May 3, 2026

0.1 2.5 Machine maintenance

States: R (running), B (broken) at the start of each week.

Terminal: $J_5(R) = J_5(B) = 0$

Immediate expected net profits and transitions:

- R + Maintain: $0.6(100-20) + 0.4(0-20) = 40$, next state: R w.p. 0.6, B w.p. 0.4
- R + No maintain: $0.3(100) + 0.7(0) = 30$, next state: R w.p. 0.3, B w.p. 0.7
- B + Repair: $0.6(100-40) + 0.4(0-40) = 20$, next state: R w.p. 0.6, B w.p. 0.4
- B + Replace: $1(100-150) = -50$, next state: R w.p. 1

Week 4:

$J_4(R)$: Maintain $\to 40$, No maintain $\to 30 \Rightarrow J_4(R) = \max(40, 30) = \mathbf{40}$ (Maintain)

$J_4(B)$: Repair $\to 20$, Replace $\to -50 \Rightarrow J_4(B) = \max(20, -50) = \mathbf{20}$ (Repair)

Week 3:

$J_3(R)$: Maintain $\to 40 + 0.6(40) + 0.4(20) = 72$, No maintain $\to 30 + 0.3(40) + 0.7(20) = 56$ $\Rightarrow J_3(R) = \max(72, 56) = \mathbf{72}$ (Maintain)

$J_3(B)$: Repair $\to 20 + 0.6(40) + 0.4(20) = 52$, Replace $\to -50 + 40 = -10 \Rightarrow J_3(B) = \max(52, -10) = \mathbf{52}$ (Repair)

Week 2:

$J_2(R)$: Maintain $\to 40 + 0.6(72) + 0.4(52) = 104$, No maintain $\to 30 + 0.3(72) + 0.7(52) = 88$ $\Rightarrow J_2(R) = \max(104, 88) = \mathbf{104}$ (Maintain)

$J_2(B)$: Repair $\to 20 + 0.6(72) + 0.4(52) = 84$, Replace $\to -50 + 72 = 22 \Rightarrow J_2(B) = \max(84, 22) = \mathbf{84}$ (Repair)

Week 1 (new machine, guaranteed to run $\Rightarrow$ earns \$100, ends in state R, no maintain needed):

$$J_1 = 100 + J_2(R) = 100 + 104 = \boxed{204}$$